



Install NX-OS software and RCFs on Cisco Nexus 9336C-FX2 cluster switches

ONTAP Systems Switches

NetApp
July 22, 2022

This PDF was generated from <https://docs.netapp.com/us-en/ontap-systems-switches/switch-cisco-9336c-fx2/install-nxos-overview.html> on July 22, 2022. Always check docs.netapp.com for the latest.

Table of Contents

- Install NX-OS software and RCFs on Cisco Nexus 9336C-FX2 cluster switches 1
 - Install NX-OS software and RCFs on Cisco Nexus 9336C-FX2 cluster switches 1
 - Install the NX-OS software 6
 - Install the Reference Configuration File (RCF) 13

Install NX-OS software and RCFs on Cisco Nexus 9336C-FX2 cluster switches

Install NX-OS software and RCFs on Cisco Nexus 9336C-FX2 cluster switches

The Cisco NX-OS software and reference configuration files (RCFs) must be installed on Cisco Nexus 9336C-FX2 cluster switches.

The following conditions must exist before you install the NX-OS software and Reference Configurations Files (RCFs) on the cluster switch:

- The cluster must be fully functioning (there should be no errors in the logs or similar issues).
- You must have checked or set your desired boot configuration in the RCF to reflect the desired boot images if you are installing only NX-OS and keeping your current RCF version.
- If you need to change the boot configuration to reflect the current boot images, you must do so before reapplying the RCF so that the correct version is instantiated on future reboots.
- You must have a console connection to the switch, required when installing the RCF.
- You must have consulted the switch compatibility table on the [Cisco Ethernet switch](#) page for the supported ONTAP, NX-OS, and RCF versions.
- There can be command dependencies between the command syntax in the RCF and that found in versions of NX-OS.
- You must have referred to the appropriate software and upgrade guides available on the Cisco web site for complete documentation on the Cisco switch upgrade and downgrade procedures on **Cisco Nexus 9000 Series Switches**.

[Cisco Nexus 9000 Series Switches](#)

- You must have the current RCF.



Before installing a new switch software version and RCFs, you must erase the switch settings and perform basic configuration. You must be connected to the switch using the serial console. This task resets the configuration of the management network.

The examples in this procedure use two nodes. These nodes use two 10GbE cluster interconnect ports e0a and e0b.

See the [Hardware Universe](#) to verify the correct cluster ports on your platforms.



The command outputs might vary depending on different releases of ONTAP.

The examples in this procedure use the following switch and node nomenclature:

- The names of the two Cisco switches are cs1 and cs2.
- The node names are cluster1-01 and cluster1-02.
- The cluster LIF names are cluster1-01_clus1 and cluster1-01_clus2 for cluster1-01 and cluster1-02_clus1 and cluster1-02_clus2 for cluster1-02.

- The `cluster1::*>` prompt indicates the name of the cluster.



The procedure requires the use of both ONTAP commands and Cisco Nexus 9000 Series Switches commands; ONTAP commands are used unless otherwise indicated.

1. If AutoSupport is enabled on this cluster, suppress automatic case creation by invoking an AutoSupport message: `system node autosupport invoke -node * -type all -message MAINT=x h`

where x is the duration of the maintenance window in hours.



The AutoSupport message notifies technical support of this maintenance task so that automatic case creation is suppressed during the maintenance window.

2. Change the privilege level to advanced, entering **y** when prompted to continue:

```
set -privilege advanced
```

The advanced prompt (`*>`) appears.

3. Display how many cluster interconnect interfaces are configured in each node for each cluster interconnect switch: `network device-discovery show -protocol cdp`

```
cluster1::*> network device-discovery show -protocol cdp
```

| Node/ Protocol | Local Port | Discovered Device (LLDP: ChassisID) | Interface | Platform |
|-------------------|---------------|--|-----------|----------|
| cluster1-02/cdp | e0a | cs1 | Eth1/2 | N9K- |
| C9336C | e0b | cs2 | Eth1/2 | N9K- |
| C9336C | | | | |
| cluster1-01/cdp | e0a | cs1 | Eth1/1 | N9K- |
| C9336C | e0b | cs2 | Eth1/1 | N9K- |
| C9336C | | | | |

4 entries were displayed.

4. Check the administrative or operational status of each cluster interface.

- a. Display the network port attributes: `network port show -ipspace Cluster`

```
cluster1::*> network port show -ipspace Cluster
```

Node: cluster1-02

| Port | IPspace | Broadcast Domain | Link | MTU | Speed(Mbps) | Health |
|------|---------|------------------|------|------|-------------|---------|
| | | | | | Admin/Oper | Status |
| e0a | Cluster | Cluster | up | 9000 | auto/10000 | healthy |
| e0b | Cluster | Cluster | up | 9000 | auto/10000 | healthy |

Node: cluster1-01

| Port | IPspace | Broadcast Domain | Link | MTU | Speed(Mbps) | Health |
|------|---------|------------------|------|------|-------------|---------|
| | | | | | Admin/Oper | Status |
| e0a | Cluster | Cluster | up | 9000 | auto/10000 | healthy |
| e0b | Cluster | Cluster | up | 9000 | auto/10000 | healthy |

4 entries were displayed.

b. Display information about the LIFs: network interface show -vserver Cluster

```
cluster1::*> network interface show -vserver Cluster
```

| Current Is | | Logical | Status | Network | Current |
|-------------|------|-------------------|------------|-------------------|---------|
| Vserver | Port | Interface | Admin/Oper | Address/Mask | Node |
| Home | | | | | |
| Cluster | | | | | |
| | | cluster1-01_clus1 | up/up | 169.254.209.69/16 | |
| cluster1-01 | | e0a true | | | |
| | | cluster1-01_clus2 | up/up | 169.254.49.125/16 | |
| cluster1-01 | | e0b true | | | |
| | | cluster1-02_clus1 | up/up | 169.254.47.194/16 | |
| cluster1-02 | | e0a true | | | |
| | | cluster1-02_clus2 | up/up | 169.254.19.183/16 | |
| cluster1-02 | | e0b true | | | |

4 entries were displayed.

5. Ping the remote cluster LIFs: cluster ping-cluster -node node-name

```

cluster1::*> cluster ping-cluster -node cluster1-02
Host is cluster1-02
Getting addresses from network interface table...
Cluster cluster1-01_clus1 169.254.209.69 cluster1-01 e0a
Cluster cluster1-01_clus2 169.254.49.125 cluster1-01 e0b
Cluster cluster1-02_clus1 169.254.47.194 cluster1-02 e0a
Cluster cluster1-02_clus2 169.254.19.183 cluster1-02 e0b
Local = 169.254.47.194 169.254.19.183
Remote = 169.254.209.69 169.254.49.125
Cluster Vserver Id = 4294967293
Ping status:

Basic connectivity succeeds on 4 path(s)
Basic connectivity fails on 0 path(s)

Detected 9000 byte MTU on 4 path(s):
    Local 169.254.19.183 to Remote 169.254.209.69
    Local 169.254.19.183 to Remote 169.254.49.125
    Local 169.254.47.194 to Remote 169.254.209.69
    Local 169.254.47.194 to Remote 169.254.49.125
Larger than PMTU communication succeeds on 4 path(s)
RPC status:
2 paths up, 0 paths down (tcp check)
2 paths up, 0 paths down (udp check)

```

6. Verify that the auto-revert command is enabled on all cluster LIFs: network interface show -vserver Cluster -fields auto-revert

```

cluster1::*> network interface show -vserver Cluster -fields auto-revert

```

| Vserver | Logical Interface | Auto-revert |
|---------|-------------------|-------------|
| Cluster | cluster1-01_clus1 | true |
| | cluster1-01_clus2 | true |
| | cluster1-02_clus1 | true |
| | cluster1-02_clus2 | true |

4 entries were displayed.

7. For ONTAP 9.8 and later, enable the Ethernet switch health monitor log collection feature for collecting switch-related log files, using the commands: system switch ethernet log setup-password and system switch ethernet log enable-collection

```

cluster1::*> system switch ethernet log setup-password
Enter the switch name: <return>
The switch name entered is not recognized.
Choose from the following list:
cs1
cs2

cluster1::*> system switch ethernet log setup-password

Enter the switch name: cs1
RSA key fingerprint is e5:8b:c6:dc:e2:18:18:09:36:63:d9:63:dd:03:d9:cc
Do you want to continue? {y|n}::[n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster1::*> system switch ethernet log setup-password

Enter the switch name: cs2
RSA key fingerprint is 57:49:86:a1:b9:80:6a:61:9a:86:8e:3c:e3:b7:1f:b1
Do you want to continue? {y|n}:: [n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster1::*> system switch ethernet log enable-collection

Do you want to enable cluster log collection for all nodes in the
cluster?
{y|n}: [n] y

Enabling cluster switch log collection.

cluster1::*>

```



If any of these commands return an error, contact NetApp support.

8. For ONTAP releases 9.5P16, 9.6P12, and 9.7P10 and later patch releases, enable the Ethernet switch health monitor log collection feature for collecting switch-related log files, using the commands: `system cluster-switch log setup-password` and `system cluster-switch log enable-collection`

```

cluster1::*> system cluster-switch log setup-password
Enter the switch name: <return>
The switch name entered is not recognized.
Choose from the following list:
cs1
cs2

cluster1::*> system cluster-switch log setup-password

Enter the switch name: cs1
RSA key fingerprint is e5:8b:c6:dc:e2:18:18:09:36:63:d9:63:dd:03:d9:cc
Do you want to continue? {y|n}::[n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster1::*> system cluster-switch log setup-password

Enter the switch name: cs2
RSA key fingerprint is 57:49:86:a1:b9:80:6a:61:9a:86:8e:3c:e3:b7:1f:b1
Do you want to continue? {y|n}:: [n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster1::*> system cluster-switch log enable-collection

Do you want to enable cluster log collection for all nodes in the
cluster?
{y|n}: [n] y

Enabling cluster switch log collection.

cluster1::*>

```



If any of these commands return an error, contact NetApp support.

Install the NX-OS software

You can use this procedure to install the NX-OS software on the Nexus 9336C-FX2 cluster switch.

Steps

1. Connect the cluster switch to the management network.

2. Use the ping command to verify connectivity to the server hosting the NX-OS software and the RCF.

This example verifies that the switch can reach the server at IP address 172.19.2.1:

```
cs2# ping 172.19.2.1
Pinging 172.19.2.1 with 0 bytes of data:

Reply From 172.19.2.1: icmp_seq = 0. time= 5910 usec.
```

3. Copy the NX-OS software and EPLD images to the Nexus 9336C-FX2 switch.

```
cs2# copy sftp: bootflash: vrf management
Enter source filename: /code/nxos.9.3.5.bin
Enter hostname for the sftp server: 172.19.2.1
Enter username: user1

Outbound-ReKey for 172.19.2.1:22
Inbound-ReKey for 172.19.2.1:22
user1@172.19.2.1's password:
sftp> progress
Progress meter enabled
sftp> get /code/nxos.9.3.5.bin /bootflash/nxos.9.3.5.bin
/code/nxos.9.3.5.bin 100% 1261MB 9.3MB/s 02:15
sftp> exit
Copy complete, now saving to disk (please wait)...
Copy complete.

cs2# copy sftp: bootflash: vrf management

Enter source filename: /code/n9000-epld.9.3.5.img
Enter hostname for the sftp server: 172.19.2.1
Enter username: user1

Outbound-ReKey for 172.19.2.1:22
Inbound-ReKey for 172.19.2.1:22
user1@172.19.2.1's password:
sftp> progress
Progress meter enabled
sftp> get /code/n9000-epld.9.3.5.img /bootflash/n9000-epld.9.3.5.img
/code/n9000-epld.9.3.5.img 100% 161MB 9.5MB/s 00:16
sftp> exit
Copy complete, now saving to disk (please wait)...
Copy complete.
```

4. Verify the running version of the NX-OS software:

```
cs2# show version
Cisco Nexus Operating System (NX-OS) Software
TAC support: http://www.cisco.com/tac
Copyright (C) 2002-2020, Cisco and/or its affiliates.
All rights reserved.
The copyrights to certain works contained in this software are
owned by other third parties and used and distributed under their own
licenses, such as open source. This software is provided "as is," and
unless
otherwise stated, there is no warranty, express or implied, including
but not
limited to warranties of merchantability and fitness for a particular
purpose.
Certain components of this software are licensed under
the GNU General Public License (GPL) version 2.0 or
GNU General Public License (GPL) version 3.0 or the GNU
Lesser General Public License (LGPL) Version 2.1 or
Lesser General Public License (LGPL) Version 2.0.
A copy of each such license is available at
http://www.opensource.org/licenses/gpl-2.0.php and
http://opensource.org/licenses/gpl-3.0.html and
http://www.opensource.org/licenses/lgpl-2.1.php and
http://www.gnu.org/licenses/old-licenses/library.txt.

Software
  BIOS: version 08.38
  NXOS: version 9.3(4)
  BIOS compile time: 05/29/2020
  NXOS image file is: bootflash:///nxos.9.3.4.bin
  NXOS compile time: 4/28/2020 21:00:00 [04/29/2020 02:28:31]

Hardware
  cisco Nexus9000 C9336C-FX2 Chassis
  Intel(R) Xeon(R) CPU E5-2403 v2 @ 1.80GHz with 8154432 kB of memory.
  Processor Board ID FOC20291J6K

  Device name: cs2
  bootflash: 53298520 kB
  Kernel uptime is 0 day(s), 0 hour(s), 3 minute(s), 42 second(s)

  Last reset at 157524 usecs after Mon Nov 2 18:32:06 2020
  Reason: Reset Requested by CLI command reload
```

System version: 9.3(4)

Service:

plugin

Core Plugin, Ethernet Plugin

Active Package(s):

cs2#

5. Install the NX-OS image.

Installing the image file causes it to be loaded every time the switch is rebooted.

```
cs2# install all nxos bootflash:nxos.9.3.5.bin
```

Installer will perform compatibility check first. Please wait.

Installer is forced disruptive

Verifying image bootflash:/nxos.9.3.5.bin for boot variable "nxos".

[#####] 100% -- SUCCESS

Verifying image type.

[#####] 100% -- SUCCESS

Preparing "nxos" version info using image bootflash:/nxos.9.3.5.bin.

[#####] 100% -- SUCCESS

Preparing "bios" version info using image bootflash:/nxos.9.3.5.bin.

[#####] 100% -- SUCCESS

Performing module support checks.

[#####] 100% -- SUCCESS

Notifying services about system upgrade.

[#####] 100% -- SUCCESS

Compatibility check is done:

| Module | bootable | Impact | Install-type | Reason |
|--------|----------|------------|--------------|--------------------------------|
| 1 | yes | disruptive | reset | default upgrade is not hitless |

Images will be upgraded according to following table:

| Module | Image | Running-Version(pri:alt) | New-Version |
|--------------------|-------|---------------------------------------|-------------|
| Upg-Required | | | |
| ----- | ----- | ----- | ----- |
| 1 | nxos | 9.3(4) | 9.3(5) |
| yes | | | |
| 1 | bios | v08.37(01/28/2020):v08.23(09/23/2015) | |
| v08.38(05/29/2020) | | yes | |

Switch will be reloaded for disruptive upgrade.

Do you want to continue with the installation (y/n)? [n] y

Install is in progress, please wait.

Performing runtime checks.

[#####] 100% -- SUCCESS

Setting boot variables.

[#####] 100% -- SUCCESS

Performing configuration copy.

[#####] 100% -- SUCCESS

Module 1: Refreshing compact flash and upgrading bios/loader/bootrom.

Warning: please do not remove or power off the module at this time.

[#####] 100% -- SUCCESS

Finishing the upgrade, switch will reboot in 10 seconds.

6. Verify the new version of NX-OS software after the switch has rebooted: show version

```
cs2# show version
```

Cisco Nexus Operating System (NX-OS) Software

TAC support: <http://www.cisco.com/tac>

Copyright (C) 2002-2020, Cisco and/or its affiliates.

All rights reserved.

The copyrights to certain works contained in this software are owned by other third parties and used and distributed under their own licenses, such as open source. This software is provided "as is," and

unless
otherwise stated, there is no warranty, express or implied, including
but not
limited to warranties of merchantability and fitness for a particular
purpose.

Certain components of this software are licensed under
the GNU General Public License (GPL) version 2.0 or
GNU General Public License (GPL) version 3.0 or the GNU
Lesser General Public License (LGPL) Version 2.1 or
Lesser General Public License (LGPL) Version 2.0.

A copy of each such license is available at
<http://www.opensource.org/licenses/gpl-2.0.php> and
<http://opensource.org/licenses/gpl-3.0.html> and
<http://www.opensource.org/licenses/lgpl-2.1.php> and
<http://www.gnu.org/licenses/old-licenses/library.txt>.

Software

BIOS: version 05.33
NXOS: version 9.3(5)
BIOS compile time: 09/08/2018
NXOS image file is: bootflash:///nxos.9.3.5.bin
NXOS compile time: 11/4/2018 21:00:00 [11/05/2018 06:11:06]

Hardware

cisco Nexus9000 C9336C-FX2 Chassis
Intel(R) Xeon(R) CPU E5-2403 v2 @ 1.80GHz with 8154432 kB of memory.
Processor Board ID FOC20291J6K

Device name: cs2
bootflash: 53298520 kB

Kernel uptime is 0 day(s), 0 hour(s), 3 minute(s), 42 second(s)

Last reset at 277524 usecs after Mon Nov 2 22:45:12 2020

Reason: Reset due to upgrade
System version: 9.3(4)
Service:

plugin

Core Plugin, Ethernet Plugin

Active Package(s):

7. Upgrade the EPLD image and reboot the switch.

```
cs2# show version module 1 epld
```

| EPLD Device | Version |
|-------------|---------|
| MI FPGA | 0x7 |
| IO FPGA | 0x17 |
| MI FPGA2 | 0x2 |
| GEM FPGA | 0x2 |
| GEM FPGA | 0x2 |
| GEM FPGA | 0x2 |
| GEM FPGA | 0x2 |

```
cs2# install epld bootflash:n9000-epld.9.3.5.img module 1
```

Compatibility check:

| Module | Type | Upgradable | Impact | Reason |
|--------|------|------------|------------|-------------------|
| 1 | SUP | Yes | disruptive | Module Upgradable |

Retrieving EPLD versions.... Please wait.

Images will be upgraded according to following table:

| Module | Type | EPLD | Running-Version | New-Version | Upg-Required |
|--------|------|----------|-----------------|-------------|--------------|
| 1 | SUP | MI FPGA | 0x07 | 0x07 | No |
| 1 | SUP | IO FPGA | 0x17 | 0x19 | Yes |
| 1 | SUP | MI FPGA2 | 0x02 | 0x02 | No |

The above modules require upgrade.

The switch will be reloaded at the end of the upgrade

Do you want to continue (y/n) ? [n] y

Proceeding to upgrade Modules.

Starting Module 1 EPLD Upgrade

Module 1 : IO FPGA [Programming] : 100.00% (64 of 64 sectors)

Module 1 EPLD upgrade is successful.

| Module | Type | Upgrade-Result |
|--------|------|----------------|
| 1 | SUP | Success |

EPLDs upgraded.

Module 1 EPLD upgrade is successful.

8. After the switch reboot, log in again and verify that the new version of EPLD loaded successfully.

```
cs2# show version module 1 epld
```

| EPLD Device | | Version |
|-------------|-------|---------|
| ----- | | |
| MI | FPGA | 0x7 |
| IO | FPGA | 0x19 |
| MI | FPGA2 | 0x2 |
| GEM | FPGA | 0x2 |
| GEM | FPGA | 0x2 |
| GEM | FPGA | 0x2 |
| GEM | FPGA | 0x2 |

Install the Reference Configuration File (RCF)

You can install the RCF after setting up the Nexus 9336C-FX2 switch for the first time. You can also use this procedure to upgrade your RCF version.

The examples in this procedure use the following switch and node nomenclature:

- The names of the two Cisco switches are cs1 and cs2.
- The node names are cluster1-01, cluster1-02, cluster1-03, and cluster1-04.
- The cluster LIF names are cluster1-01_clus1, cluster1-01_clus2, cluster1-02_clus1, cluster1-02_clus2, cluster1-03_clus1, cluster1-03_clus2, cluster1-04_clus1, and cluster1-04_clus2.
- The `cluster1::*>` prompt indicates the name of the cluster.



- The procedure requires the use of both ONTAP commands and Cisco Nexus 9000 Series Switches commands; ONTAP commands are used unless otherwise indicated.
- Before you perform this procedure make sure that you have a current backup of the switch configuration.

Steps

1. Display the cluster ports on each node that are connected to the cluster switches: `network device-discovery show`

```

cluster1::*> network device-discovery show
Node/      Local  Discovered
Protocol   Port   Device (LLDP: ChassisID)  Interface      Platform
-----
cluster1-01/cdp
           e0a    cs1                      Ethernet1/7     N9K-
C9336C
           e0d    cs2                      Ethernet1/7     N9K-
C9336C
cluster1-02/cdp
           e0a    cs1                      Ethernet1/8     N9K-
C9336C
           e0d    cs2                      Ethernet1/8     N9K-
C9336C
cluster1-03/cdp
           e0a    cs1                      Ethernet1/1/1   N9K-
C9336C
           e0b    cs2                      Ethernet1/1/1   N9K-
C9336C
cluster1-04/cdp
           e0a    cs1                      Ethernet1/1/2   N9K-
C9336C
           e0b    cs2                      Ethernet1/1/2   N9K-
C9336C
cluster1::*>

```

2. Check the administrative and operational status of each cluster port.

- a. Verify that all the cluster ports are up with a healthy status: `network port show -role cluster`

```

cluster1::*> network port show -role cluster

Node: cluster1-01

Ignore

Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
Status
-----
e0a       Cluster      Cluster      up   9000  auto/100000
healthy false
e0d       Cluster      Cluster      up   9000  auto/100000
healthy false

```


Node: cluster1-02

Ignore

| | | | | | | Speed(Mbps) | Health |
|---------|---------|-----------|--------|------|------|-------------|--------|
| Health | | | | | | | |
| Port | IPspace | Broadcast | Domain | Link | MTU | Admin/Oper | Status |
| Status | | | | | | | |
| ----- | ----- | ----- | | ---- | ---- | ----- | |
| ----- | ----- | | | | | | |
| e0a | Cluster | Cluster | | up | 9000 | auto/100000 | |
| healthy | false | | | | | | |
| e0d | Cluster | Cluster | | up | 9000 | auto/100000 | |
| healthy | false | | | | | | |

8 entries were displayed.

Node: cluster1-03

Ignore

| | | | | | | Speed(Mbps) | Health |
|--------|---------|-----------|--------|------|------|-------------|---------|
| Health | | | | | | | |
| Port | IPspace | Broadcast | Domain | Link | MTU | Admin/Oper | Status |
| Status | | | | | | | |
| ----- | ----- | ----- | | ---- | ---- | ----- | |
| ----- | ----- | | | | | | |
| e0a | Cluster | Cluster | | up | 9000 | auto/10000 | healthy |
| false | | | | | | | |
| e0b | Cluster | Cluster | | up | 9000 | auto/10000 | healthy |
| false | | | | | | | |

Node: cluster1-04

Ignore

| | | | | | | Speed(Mbps) | Health |
|--------|---------|-----------|--------|------|------|-------------|---------|
| Health | | | | | | | |
| Port | IPspace | Broadcast | Domain | Link | MTU | Admin/Oper | Status |
| Status | | | | | | | |
| ----- | ----- | ----- | | ---- | ---- | ----- | |
| ----- | ----- | | | | | | |
| e0a | Cluster | Cluster | | up | 9000 | auto/10000 | healthy |
| false | | | | | | | |
| e0b | Cluster | Cluster | | up | 9000 | auto/10000 | healthy |
| false | | | | | | | |

cluster1::*>

b. Verify that all the cluster interfaces (LIFs) are on the home port: `network interface show -role`

cluster

```
cluster1::*> network interface show -role cluster
```

| | | Logical | Status | Network | Current |
|------------|-----------|---------------------------|------------|----------------|-----------|
| Current Is | | | | | |
| Vserver | Interface | | Admin/Oper | Address/Mask | Node |
| Port | Home | | | | |
| ----- | | | | | |
| ----- | | | | | |
| Cluster | | | | | |
| 01 | e0a | cluster1-01_clus1 true | up/up | 169.254.3.4/23 | cluster1- |
| 01 | e0d | cluster1-01_clus2 true | up/up | 169.254.3.5/23 | cluster1- |
| 02 | e0a | cluster1-02_clus1 true | up/up | 169.254.3.8/23 | cluster1- |
| 02 | e0d | cluster1-02_clus2 true | up/up | 169.254.3.9/23 | cluster1- |
| 03 | e0a | cluster1-03_clus1 true | up/up | 169.254.1.3/23 | cluster1- |
| 03 | e0b | cluster1-03_clus2 true | up/up | 169.254.1.1/23 | cluster1- |
| 04 | e0a | cluster1-04_clus1 true | up/up | 169.254.1.6/23 | cluster1- |
| 04 | e0b | cluster1-04_clus2 true | up/up | 169.254.1.7/23 | cluster1- |

8 entries were displayed.
cluster1::*>

- c. Verify that the cluster displays information for both cluster switches: `system cluster-switch show -is-monitoring-enabled-operational true`

```
cluster1::*> system cluster-switch show -is-monitoring-enabled
-operational true
```

| Switch | Type | Address | Model |
|---|-----------------|---------------|------------|
| cs1 | cluster-network | 10.233.205.90 | N9K-C9336C |
| Serial Number: FOCXXXXXXGD Is Monitored: true Reason: None Software Version: Cisco Nexus Operating System (NX-OS) Software, Version 9.3(5) Version Source: CDP | | | |
| cs2 | cluster-network | 10.233.205.91 | N9K-C9336C |
| Serial Number: FOCXXXXXXGS Is Monitored: true Reason: None Software Version: Cisco Nexus Operating System (NX-OS) Software, Version 9.3(5) Version Source: CDP | | | |

```
cluster1::*>
```

3. Disable auto-revert on the cluster LIFs.

```
cluster1::*> network interface modify -vserver Cluster -lif \* -auto
-revert false
```

4. On cluster switch cs2, shut down the ports connected to the cluster ports of the nodes.

```
cs2(config)# interface eth1/1/1-2,eth1/7-8
cs2(config-if-range)# shutdown
```

5. Verify that the cluster LIFs have migrated to the ports hosted on cluster switch cs1. This might take a few seconds.

```
network interface show -role cluster
```

```
cluster1::*> network interface show -role cluster
```

| | Logical | Status | Network | Current |
|------------|-------------------|------------|----------------|-------------|
| Current Is | | | | |
| Vserver | Interface | Admin/Oper | Address/Mask | Node |
| Port | Home | | | |
| ----- | | | | |
| Cluster | | | | |
| | cluster1-01_clus1 | up/up | 169.254.3.4/23 | cluster1-01 |
| e0a | true | | | |
| | cluster1-01_clus2 | up/up | 169.254.3.5/23 | cluster1-01 |
| e0a | false | | | |
| | cluster1-02_clus1 | up/up | 169.254.3.8/23 | cluster1-02 |
| e0a | true | | | |
| | cluster1-02_clus2 | up/up | 169.254.3.9/23 | cluster1-02 |
| e0a | false | | | |
| | cluster1-03_clus1 | up/up | 169.254.1.3/23 | cluster1-03 |
| e0a | true | | | |
| | cluster1-03_clus2 | up/up | 169.254.1.1/23 | cluster1-03 |
| e0a | false | | | |
| | cluster1-04_clus1 | up/up | 169.254.1.6/23 | cluster1-04 |
| e0a | true | | | |
| | cluster1-04_clus2 | up/up | 169.254.1.7/23 | cluster1-04 |
| e0a | false | | | |

8 entries were displayed.
cluster1::*>

6. Verify that the cluster is healthy: `cluster show`

```
cluster1::*> cluster show
```

| Node | Health | Eligibility | Epsilon |
|-------------|--------|-------------|---------|
| ----- | ----- | ----- | ----- |
| cluster1-01 | true | true | false |
| cluster1-02 | true | true | false |
| cluster1-03 | true | true | true |
| cluster1-04 | true | true | false |

4 entries were displayed.
cluster1::*>

7. If you have not already done so, save the current switch configuration by copying the output of the following command to a log file:

```
show running-config
```

8. Clean the configuration on switch cs2 and perform a basic setup.



When updating or applying a new RCF, you must erase the switch settings and perform basic configuration. You must be connected to the switch using the serial console to erase switch settings.

a. Clean the configuration. This step requires a console connection to the switch.

```
cs2# write erase
Warning: This command will erase the startup-configuration.
Do you wish to proceed anyway? (y/n)  [n] y
cs2# reload
This command will reboot the system. (y/n)?  [n] y
cs2#
```

b. Perform a basic setup of the switch.

9. Copy the RCF to the bootflash of switch cs2 using one of the following transfer protocols: FTP, TFTP, SFTP, or SCP. For more information on Cisco commands, see the appropriate guide in the [Cisco Nexus 9000 Series NX-OS Command Reference](#) guides.

This example shows TFTP being used to copy an RCF to the bootflash on switch cs2:

```
cs2# copy tftp: bootflash: vrf management
Enter source filename: Nexus_9336C_RCF_v1.6-Cluster-HA-Breakout.txt
Enter hostname for the tftp server: 172.22.201.50
Trying to connect to tftp server.....Connection to Server Established.
TFTP get operation was successful
Copy complete, now saving to disk (please wait)...
```

10. Apply the RCF previously downloaded to the bootflash.

For more information on Cisco commands, see the appropriate guide in the [Cisco Nexus 9000 Series NX-OS Command Reference](#) guides.

This example shows the RCF file `Nexus_9336C_RCF_v1.6-Cluster-HA-Breakout.txt` being installed on switch cs2:

```
cs2# copy Nexus_9336C_RCF_v1.6-Cluster-HA-Breakout.txt running-config
echo-commands
```

11. Examine the banner output from the `show banner motd` command. You must read and follow these instructions to ensure the proper configuration and operation of the switch.

```
cs2# show banner motd
```

```
*****
*****
* NetApp Reference Configuration File (RCF)
*
* Switch    : Nexus N9K-C9336C-FX2
* Filename  : Nexus_9336C_RCF_v1.6-Cluster-HA-Breakout.txt
* Date      : 10-23-2020
* Version   : v1.6
*
* Port Usage:
* Ports 1- 3: Breakout mode (4x10G) Intra-Cluster Ports, int e1/1/1-4,
e1/2/1-4
, e1/3/1-4
* Ports 4- 6: Breakout mode (4x25G) Intra-Cluster/HA Ports, int e1/4/1-
4, e1/5/
1-4, e1/6/1-4
* Ports 7-34: 40/100GbE Intra-Cluster/HA Ports, int e1/7-34
* Ports 35-36: Intra-Cluster ISL Ports, int e1/35-36
*
* Dynamic breakout commands:
* 10G: interface breakout module 1 port <range> map 10g-4x
* 25G: interface breakout module 1 port <range> map 25g-4x
*
* Undo breakout commands and return interfaces to 40/100G configuration
in confi
g mode:
* no interface breakout module 1 port <range> map 10g-4x
* no interface breakout module 1 port <range> map 25g-4x
* interface Ethernet <interfaces taken out of breakout mode>
* inherit port-profile 40-100G
* priority-flow-control mode auto
* service-policy input HA
* exit
*
*****
*****
```

12. Verify that the RCF file is the correct newer version: `show running-config`

When you check the output to verify you have the correct RCF, make sure that the following information is correct:

- The RCF banner

- The node and port settings
- Customizations The output varies according to your site configuration. Check the port settings and refer to the release notes for any changes specific to the RCF that you have installed.

13. After you verify the RCF versions and switch settings are correct, copy the running-config file to the startup-config file.

For more information on Cisco commands, see the appropriate guide in the [Cisco Nexus 9000 Series NX-OS Command Reference](#) guides.

```
cs2# copy running-config startup-config
[#####] 100% Copy complete
```

14. Reboot switch cs2. You can ignore the “cluster ports down” events reported on the nodes while the switch reboots.

```
cs2# reload
This command will reboot the system. (y/n)? [n] y
```

15. Verify the health of cluster ports on the cluster.

- Verify that e0d ports are up and healthy across all nodes in the cluster: `network port show -role cluster`

```
cluster1::*> network port show -role cluster

Node: cluster1-01

Ignore
Speed(Mbps) Health
Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
Status
-----
e0a      Cluster      Cluster      up    9000  auto/10000 healthy
false
e0b      Cluster      Cluster      up    9000  auto/10000 healthy
false

Node: cluster1-02

Ignore
Speed(Mbps) Health
Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
```

```

Status
-----
e0a      Cluster      Cluster      up    9000    auto/10000 healthy
false
e0b      Cluster      Cluster      up    9000    auto/10000 healthy
false

Node: cluster1-03

Ignore

Speed(Mbps) Health
Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
Status
-----
e0a      Cluster      Cluster      up    9000    auto/100000
healthy false
e0d      Cluster      Cluster      up    9000    auto/100000
healthy false

Node: cluster1-04

Ignore

Speed(Mbps) Health
Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
Status
-----
e0a      Cluster      Cluster      up    9000    auto/100000
healthy false
e0d      Cluster      Cluster      up    9000    auto/100000
healthy false
8 entries were displayed.

```

- b. Verify the switch health from the cluster (this might not show switch cs2, since LIFs are not homed on e0d).

```

cluster1::*> network device-discovery show -protocol cdp
Node/      Local  Discovered
Protocol   Port   Device (LLDP: ChassisID)  Interface
Platform
-----

```



```

cluster1-01/cdp
    e0a    cs1                                Ethernet1/7    N9K-
C9336C
    e0d    cs2                                Ethernet1/7    N9K-
C9336C
cluster01-2/cdp
    e0a    cs1                                Ethernet1/8    N9K-
C9336C
    e0d    cs2                                Ethernet1/8    N9K-
C9336C
cluster01-3/cdp
    e0a    cs1                                Ethernet1/1/1  N9K-
C9336C
    e0b    cs2                                Ethernet1/1/1  N9K-
C9336C
cluster1-04/cdp
    e0a    cs1                                Ethernet1/1/2  N9K-
C9336C
    e0b    cs2                                Ethernet1/1/2  N9K-
C9336C

cluster1::*> system cluster-switch show -is-monitoring-enabled
-operational true
Switch                                     Type                Address             Model
-----
cs1                                       cluster-network     10.233.205.90      NX9-
C9336C
    Serial Number: FOCXXXXXXGD
    Is Monitored: true
    Reason: None
    Software Version: Cisco Nexus Operating System (NX-OS) Software,
Version
                        9.3(5)
    Version Source: CDP

cs2                                       cluster-network     10.233.205.91      NX9-
C9336C
    Serial Number: FOCXXXXXXGS
    Is Monitored: true
    Reason: None
    Software Version: Cisco Nexus Operating System (NX-OS) Software,
Version
                        9.3(5)
    Version Source: CDP

2 entries were displayed.

```



You might observe the following output on the cs1 switch console depending on the RCF version previously loaded on the switch:

```
2020 Nov 17 16:07:18 cs1 %$ VDC-1 %$ %STP-2-UNBLOCK_CONSIST_PORT:
Unblocking port port-channel1 on VLAN0092. Port consistency restored.
2020 Nov 17 16:07:23 cs1 %$ VDC-1 %$ %STP-2-BLOCK_PVID_PEER: Blocking
port-channel1 on VLAN0001. Inconsistent peer vlan.
2020 Nov 17 16:07:23 cs1 %$ VDC-1 %$ %STP-2-BLOCK_PVID_LOCAL: Blocking
port-channel1 on VLAN0092. Inconsistent local vlan.
```

16. On cluster switch cs1, shut down the ports connected to the cluster ports of the nodes.

The following example uses the interface example output from step 1:

```
cs1(config)# interface eth1/1/1-2,eth1/7-8
cs1(config-if-range)# shutdown
```

17. Verify that the cluster LIFs have migrated to the ports hosted on switch cs2. This might take a few seconds.
network interface show -role cluster

```

cluster1::*> network interface show -role cluster

```

| Current Is | Logical | Status | Network | Current |
|------------|-------------------|------------|----------------|-------------|
| Vserver | Interface | Admin/Oper | Address/Mask | Node |
| Port | Home | | | |
| ----- | | | | |
| Cluster | | | | |
| e0d | cluster1-01_clus1 | up/up | 169.254.3.4/23 | cluster1-01 |
| | false | | | |
| e0d | cluster1-01_clus2 | up/up | 169.254.3.5/23 | cluster1-01 |
| | true | | | |
| e0d | cluster1-02_clus1 | up/up | 169.254.3.8/23 | cluster1-02 |
| | false | | | |
| e0d | cluster1-02_clus2 | up/up | 169.254.3.9/23 | cluster1-02 |
| | true | | | |
| e0b | cluster1-03_clus1 | up/up | 169.254.1.3/23 | cluster1-03 |
| | false | | | |
| e0b | cluster1-03_clus2 | up/up | 169.254.1.1/23 | cluster1-03 |
| | true | | | |
| e0b | cluster1-04_clus1 | up/up | 169.254.1.6/23 | cluster1-04 |
| | false | | | |
| e0b | cluster1-04_clus2 | up/up | 169.254.1.7/23 | cluster1-04 |
| | true | | | |

```

8 entries were displayed.
cluster1::*>

```

18. Verify that the cluster is healthy: cluster show

```

cluster1::*> cluster show

```

| Node | Health | Eligibility | Epsilon |
|-------------|--------|-------------|---------|
| cluster1-01 | true | true | false |
| cluster1-02 | true | true | false |
| cluster1-03 | true | true | true |
| cluster1-04 | true | true | false |

```

4 entries were displayed.
cluster1::*>

```

19. Repeat Steps 7 to 14 on switch cs1.

20. Enable auto-revert on the cluster LIFs.

```
cluster1::*> network interface modify -vserver Cluster -lif \* -auto  
-revert True
```

21. Reboot switch cs1. You do this to trigger the cluster LIFs to revert to their home ports. You can ignore the “cluster ports down” events reported on the nodes while the switch reboots.

```
cs1# reload  
This command will reboot the system. (y/n)? [n] y
```

22. Verify that the switch ports connected to the cluster ports are up.

```
cs1# show interface brief \ | grep up  
.   
.   
Eth1/1/1      1      eth  access up      none      10G(D)  
--   
Eth1/1/2      1      eth  access up      none      10G(D)  
--   
Eth1/7        1      eth  trunk  up      none      100G(D)  
--   
Eth1/8        1      eth  trunk  up      none      100G(D)  
--   
.   
.
```

23. Verify that the ISL between cs1 and cs2 is functional: `show port-channel summary`

```
cs1# show port-channel summary
```

```
Flags:  D - Down          P - Up in port-channel (members)
        I - Individual    H - Hot-standby (LACP only)
        s - Suspended     r - Module-removed
        b - BFD Session Wait
        S - Switched      R - Routed
        U - Up (port-channel)
        p - Up in delay-lACP mode (member)
        M - Not in use. Min-links not met
```

```
-----
-----
Group Port-          Type      Protocol  Member Ports      Channel
-----
-----
1      Po1 (SU)      Eth      LACP      Eth1/35 (P)      Eth1/36 (P)
cs1#
```

24. Verify that the cluster LIFs have reverted to their home port: `network interface show -role cluster`

```

cluster1::*> network interface show -role cluster

```

| Current Is | Logical | Status | Network | Current |
|------------|-------------------|------------|----------------|-------------|
| Vserver | Interface | Admin/Oper | Address/Mask | Node |
| Port | Home | | | |
| ----- | | | | |
| Cluster | | | | |
| e0d | cluster1-01_clus1 | up/up | 169.254.3.4/23 | cluster1-01 |
| e0d | cluster1-01_clus2 | up/up | 169.254.3.5/23 | cluster1-01 |
| e0d | cluster1-02_clus1 | up/up | 169.254.3.8/23 | cluster1-02 |
| e0d | cluster1-02_clus2 | up/up | 169.254.3.9/23 | cluster1-02 |
| e0b | cluster1-03_clus1 | up/up | 169.254.1.3/23 | cluster1-03 |
| e0b | cluster1-03_clus2 | up/up | 169.254.1.1/23 | cluster1-03 |
| e0b | cluster1-04_clus1 | up/up | 169.254.1.6/23 | cluster1-04 |
| e0b | cluster1-04_clus2 | up/up | 169.254.1.7/23 | cluster1-04 |

```

8 entries were displayed.
cluster1::*>

```

25. Verify that the cluster is healthy: cluster show

```

cluster1::*> cluster show

```

| Node | Health | Eligibility | Epsilon |
|-------------|--------|-------------|---------|
| cluster1-01 | true | true | false |
| cluster1-02 | true | true | false |
| cluster1-03 | true | true | true |
| cluster1-04 | true | true | false |

```

4 entries were displayed.
cluster1::*>

```

26. Ping the remote cluster interfaces to verify connectivity: cluster ping-cluster -node local

```

cluster1::*> cluster ping-cluster -node local
Host is cluster1-03
Getting addresses from network interface table...
Cluster cluster1-03_clus1 169.254.1.3 cluster1-03 e0a
Cluster cluster1-03_clus2 169.254.1.1 cluster1-03 e0b
Cluster cluster1-04_clus1 169.254.1.6 cluster1-04 e0a
Cluster cluster1-04_clus2 169.254.1.7 cluster1-04 e0b
Cluster cluster1-01_clus1 169.254.3.4 cluster1-01 e0a
Cluster cluster1-01_clus2 169.254.3.5 cluster1-01 e0d
Cluster cluster1-02_clus1 169.254.3.8 cluster1-02 e0a
Cluster cluster1-02_clus2 169.254.3.9 cluster1-02 e0d
Local = 169.254.1.3 169.254.1.1
Remote = 169.254.1.6 169.254.1.7 169.254.3.4 169.254.3.5 169.254.3.8
169.254.3.9
Cluster Vserver Id = 4294967293
Ping status:
.....
Basic connectivity succeeds on 12 path(s)
Basic connectivity fails on 0 path(s)
.....
Detected 9000 byte MTU on 12 path(s):
    Local 169.254.1.3 to Remote 169.254.1.6
    Local 169.254.1.3 to Remote 169.254.1.7
    Local 169.254.1.3 to Remote 169.254.3.4
    Local 169.254.1.3 to Remote 169.254.3.5
    Local 169.254.1.3 to Remote 169.254.3.8
    Local 169.254.1.3 to Remote 169.254.3.9
    Local 169.254.1.1 to Remote 169.254.1.6
    Local 169.254.1.1 to Remote 169.254.1.7
    Local 169.254.1.1 to Remote 169.254.3.4
    Local 169.254.1.1 to Remote 169.254.3.5
    Local 169.254.1.1 to Remote 169.254.3.8
    Local 169.254.1.1 to Remote 169.254.3.9
Larger than PMTU communication succeeds on 12 path(s)
RPC status:
6 paths up, 0 paths down (tcp check)
6 paths up, 0 paths down (udp check)

```

Copyright Information

Copyright © 2022 NetApp, Inc. All rights reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means-graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system- without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP "AS IS" AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark Information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.