Migrate from a Cisco switch to a Cisco Nexus 9336C-FX2 cluster switch

ONTAP Systems Switches

NetApp
April 18, 2022

Table of Contents

Migrate from a Cisco switch to a Cisco Nexus 9336C-FX2 cluster switch

Migrate from a Cisco switch to a Cisco Nexus 9336C-FX2 cluster switch

You must be aware of certain configuration information, port connections and cabling requirements when you are replacing some older Cisco Nexus cluster switches with Cisco Nexus 9336C-FX2 cluster switches.

- The following cluster switches are supported:
  - Nexus 9336C-FX2
  - Nexus 92300YC
  - Nexus 5596UP
  - Nexus 3232C
  - Nexus 3132Q-V
- See the Hardware Universe for full details of supported ports and their configurations.
- You have configured some of the ports on Nexus 9336C-FX2 switches to run at 10 GbE or 40 GbE.
- You have planned, migrated, and documented 10 GbE and 40 GbE connectivity from nodes to Nexus 9336C-FX2 cluster switches.
- The ONTAP and NX-OS versions supported in this procedure are on the Cisco Ethernet Switches page.

How to migrate from a Cisco switch to a Cisco Nexus 9336C-FX2 cluster switch

You can migrate nondisruptively older Cisco cluster switches for an ONTAP cluster to Cisco Nexus 9336C-FX2 cluster network switches.

- The existing cluster must be properly set up and functioning.
- All cluster ports must be in the up state to ensure nondisruptive operations.
- The Nexus 9336C-FX2 cluster switches must be configured and operating under the proper version of NX-OS installed and reference configuration file (RCF) applied.
- The existing cluster network configuration must have the following:
  - A redundant and fully functional NetApp cluster using both older Cisco switches.
  - Management connectivity and console access to both the older Cisco switches and the new switches.
  - All cluster LIFs in the up state with the cluster LIfs are on their home ports.
  - ISL ports enabled and cabled between the older Cisco switches and between the new switches.

The examples in this procedure use the following switch and node nomenclature:

- The existing Cisco Nexus 5596UP cluster switches are c1 and c2.

- The new Nexus 9336C-FX2 cluster switches are cs1 and cs2.
- The nodes are node1 and node2.
- The cluster LIFs are node1_clus1 and node1_clus2 on node 1, and node2_clus1 and node2_clus2 on node 2 respectively.
- Switch c2 is replaced by switch cs2 first and then switch c1 is replaced by switch cs1.
  - A temporary ISL is built on cs2 connecting c2 to cs2.
  - Cabling between the nodes and c2 are then disconnected from c2 and reconnected to cs2.
  - Cabling between the nodes and c1 are then disconnected from c1 and reconnected to cs1.
  - The temporary ISL between c2 and cs2 is then removed.

Steps

1. If AutoSupport is enabled on this cluster, suppress automatic case creation by invoking an AutoSupport message: `system node autosupport invoke -node * -type all -message MAINT=xh`

   where x is the duration of the maintenance window in hours.

   > (i) The AutoSupport message notifies technical support of this maintenance task so that automatic case creation is suppressed during the maintenance window.

2. Change the privilege level to advanced, entering **y** when prompted to continue: `set -privilege advanced`

   The advanced prompt (*>) appears.

3. Verify that auto-revert is enabled on all cluster LIFs: `network interface show -vserver Cluster -fields auto-revert`

   ```
   cluster1::*> network interface show -vserver Cluster -fields auto-revert

             Logical
   Vserver   Interface      Auto-revert
   --------- -------------  ------------
   Cluster
             node1_clus1    true
             node1_clus2    true
             node2_clus1    true
             node2_clus2    true

   4 entries were displayed.
   ```

4. Determine the administrative or operational status for each cluster interface:

   Each port should display up for `Link` and healthy for `Health Status`.

   a. Display the network port attributes: `network port show -ipspace Cluster`

```
cluster1::*> network port show -ipspace Cluster

Node: node1

Ignore
                                                  Speed(Mbps) Health
Health
Port      IPspace        Broadcast Domain Link MTU  Admin/Oper  Status
Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e0a       Cluster        Cluster          up   9000  auto/10000 healthy
false
e0b       Cluster        Cluster          up   9000  auto/10000 healthy
false

Node: node2

Ignore
                                                  Speed(Mbps) Health
Health
Port      IPspace        Broadcast Domain Link MTU  Admin/Oper  Status
Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e0a       Cluster        Cluster          up   9000  auto/10000 healthy
false
e0b       Cluster        Cluster          up   9000  auto/10000 healthy
false

4 entries were displayed.
```

b. Display information about the logical interfaces and their designated home nodes: `network interface show -vserver Cluster`

   Each LIF should display up/up for `Status Admin/Oper` and true for `Is Home`.

```
cluster1::*> network interface show -vserver Cluster

          Logical       Status      Network               Current
Current Is
Vserver     Interface   Admin/Oper  Address/Mask          Node
Port    Home
----------- ----------- ----------- -----------------     -------------
------- ----
Cluster
          node1_clus1   up/up       169.254.209.69/16     node1
e0a       true
          node1_clus2   up/up       169.254.49.125/16     node1
e0b       true
          node2_clus1   up/up       169.254.47.194/16     node2
e0a       true
          node2_clus2   up/up       169.254.19.183/16     node2
e0b       true

4 entries were displayed.
```

5. The cluster ports on each node are connected to existing cluster switches in the following way (from the nodes' perspective) using the command: `network device-discovery show -protocol cdp`

```
cluster1::*> network device-discovery show -protocol cdp
Node/        Local  Discovered
Protocol     Port   Device (LLDP: ChassisID)  Interface         Platform
----------- ------ ------------------------- ----------------
----------------
node1        /cdp
          e0a    c1                                0/1             N5K-
C5596UP
          e0b    c2                                0/1             N5K-
C5596UP
node2        /cdp
          e0a    c1                                0/2             N5K-
C5596UP
          e0b    c2                                0/2             N5K-
C5596UP
```

6. The cluster ports and switches are connected in the following way (from the switches' perspective) using the command: `show cdp neighbors`

```
c1# show cdp neighbors
```

```
Capability Codes: R - Router, T - Trans-Bridge, B - Source-Route-Bridge
                  S - Switch, H - Host, I - IGMP, r - Repeater,
                  V - VoIP-Phone, D - Remotely-Managed-Device,
                  s - Supports-STP-Dispute


Device-ID          Local Intrfce Hldtme Capability  Platform
Port ID
node1              Eth1/1        124    H           FAS2750
e0a
node2              Eth1/2        124    H           FAS2750
e0a
c2                 Eth1/41       179    S I s       N5K-C5596UP
Eth1/41
c2                 Eth1/42       175    S I s       N5K-C5596UP
Eth1/42
c2                 Eth1/43       179    S I s       N5K-C5596UP
Eth1/43
c2                 Eth1/44       175    S I s       N5K-C5596UP
Eth1/44
c2                 Eth1/45       179    S I s       N5K-C5596UP
Eth1/45
c2                 Eth1/46       179    S I s       N5K-C5596UP
Eth1/46
c2                 Eth1/47       175    S I s       N5K-C5596UP
Eth1/47
c2                 Eth1/48       179    S I s       N5K-C5596UP
Eth1/48
Total entries displayed: 10

c2# show cdp neighbors

Capability Codes: R - Router, T - Trans-Bridge, B - Source-Route-Bridge
                  S - Switch, H - Host, I - IGMP, r - Repeater,
                  V - VoIP-Phone, D - Remotely-Managed-Device,
                  s - Supports-STP-Dispute


Device-ID          Local Intrfce Hldtme Capability  Platform
Port ID
node1              Eth1/1        124    H           FAS2750
e0b
node2              Eth1/2        124    H           FAS2750
e0b
c1                 Eth1/41       175    S I s       N5K-C5596UP
Eth1/41
```

```
c1                          Eth1/42        175    S I s        N5K-C5596UP
Eth1/42
c1                          Eth1/43        175    S I s        N5K-C5596UP
Eth1/43
c1                          Eth1/44        175    S I s        N5K-C5596UP
Eth1/44
c1                          Eth1/45        175    S I s        N5K-C5596UP
Eth1/45
c1                          Eth1/46        175    S I s        N5K-C5596UP
Eth1/46
c1                          Eth1/47        176    S I s        N5K-C5596UP
Eth1/47
c1                          Eth1/48        176    S I s        N5K-C5596UP
Eth1/48
```

7. Ensure that the cluster network has full connectivity using the command: `cluster ping-cluster -node node-name`

```
cluster1::*> cluster ping-cluster -node node2

Host is node2
Getting addresses from network interface table...
Cluster node1_clus1 169.254.209.69 node1     e0a
Cluster node1_clus2 169.254.49.125 node1     e0b
Cluster node2_clus1 169.254.47.194 node2     e0a
Cluster node2_clus2 169.254.19.183 node2     e0b
Local = 169.254.47.194 169.254.19.183
Remote = 169.254.209.69 169.254.49.125
Cluster Vserver Id = 4294967293
Ping status:
....
Basic connectivity succeeds on 4 path(s)
Basic connectivity fails on 0 path(s)
................
Detected 9000 byte MTU on 4 path(s):
    Local 169.254.19.183 to Remote 169.254.209.69
    Local 169.254.19.183 to Remote 169.254.49.125
    Local 169.254.47.194 to Remote 169.254.209.69
    Local 169.254.47.194 to Remote 169.254.49.125
Larger than PMTU communication succeeds on 4 path(s)
RPC status:
2 paths up, 0 paths down (tcp check)
2 paths up, 0 paths down (udp check)
```

8. Configure a temporary ISL on cs1 on ports e1/33-34, between c1 and cs1.

The following example shows how the new ISL is configured on c1 and cs1:

```
cs2# configure
Enter configuration commands, one per line. End with CNTL/Z.
cs2(config)# interface e1/33-34
cs2(config-if-range)# description temporary ISL between Nexus 5596UP and
Nexus 9336C
cs2(config-if-range)# no lldp transmit
cs2(config-if-range)# no lldp receive
cs2(config-if-range)# switchport mode trunk
cs2(config-if-range)# no spanning-tree bpduguard enable
cs2(config-if-range)# channel-group 101 mode active
cs2(config-if-range)# exit
cs2(config)# interface port-channel 101
cs2(config-if)# switchport mode trunk
cs2(config-if)# spanning-tree port type network
cs2(config-if)# exit
cs2(config)# exit
```

9. Remove ISL cables from ports e1/33-34 from c2 and connect the cables to ports e1/33-34 on cs2.

10. Verify that the ISL ports and port-channel are operational connecting c2 and cs2: `show port-channel summary`

The following example shows the Cisco show port-channel summary command being used to verify the ISL ports are operational on c2 and cs2:

```
c2# show port-channel summary
Flags:  D - Down         P - Up in port-channel (members)
        I - Individual   H - Hot-standby (LACP only)
        s - Suspended    r - Module-removed
        b - BFD Session Wait
        S - Switched     R - Routed
        U - Up (port-channel)
        p - Up in delay-lacp mode (member)
        M - Not in use. Min-links not met
--------------------------------------------------------------------
--------
Group Port-       Type     Protocol  Member Ports
      Channel
--------------------------------------------------------------------
--------
1     Po1(SU)     Eth      LACP      Eth1/41(P)    Eth1/42(P)
Eth1/43(P)
                                     Eth1/44(P)    Eth1/45(P)
Eth1/46(P)
                                     Eth1/47(P)    Eth1/48(P)


cs2# show port-channel summary
Flags:  D - Down         P - Up in port-channel (members)
        I - Individual   H - Hot-standby (LACP only)
        s - Suspended    r - Module-removed
        b - BFD Session Wait
        S - Switched     R - Routed
        U - Up (port-channel)
        p - Up in delay-lacp mode (member)
        M - Not in use. Min-links not met
--------------------------------------------------------------------
--------
Group Port-       Type     Protocol  Member Ports
      Channel
--------------------------------------------------------------------
--------
1     Po1(SU)     Eth      LACP      Eth1/35(P)    Eth1/36(P)
101   Po101(SU)   Eth      LACP      Eth1/41(P)    Eth1/42(P)
Eth1/43(P)
                                     Eth1/44(P)    Eth1/45(P)
Eth1/46(P)
                                     Eth1/47(P)    Eth1/48(P)
```

11. For node1, disconnect the cable from e1/1 on c2, and then connect the cable to e1/1 on cs2, using

appropriate cabling supported by Nexus 9336C-FX2.

12. For node2, disconnect the cable from e1/2 on c2, and then connect the cable to e1/2 on cs2, using appropriate cabling supported by Nexus 9336C-FX2.

13. The cluster ports on each node are now connected to cluster switches in the following way, from the nodes' perspective: `network device-discovery show -protocol cdp`

```
cluster1::*> network device-discovery show -protocol cdp

Node/         Local  Discovered
Protocol      Port   Device (LLDP: ChassisID)  Interface        Platform
-----------   ------ ------------------------- ----------------
----------------
node1         /cdp
              e0a    c1                         0/1              N5K-
C5596UP
              e0b    cs2                        0/1              N9K-
C9336C

node2         /cdp
              e0a    c1                         0/2              N5K-
C5596UP
              e0b    cs2                        0/2              N9K-
C9336C
```

14. For node1, disconnect the cable from e1/1 on c1, and then connect the cable to e1/1 on cs1, using appropriate cabling supported by Nexus 9336C-FX2.

15. For node2, disconnect the cable from e1/2 on c1, and then connect the cable to e1/2 on cs1, using appropriate cabling supported by Nexus 9336C-FX2.

16. The cluster ports on each node are now connected to cluster switches in the following way, from the nodes' perspective: `network device-discovery show -protocol cdp`

```
cluster1::*> network device-discovery show -protocol cdp
Node/          Local   Discovered
Protocol       Port    Device (LLDP: ChassisID)   Interface       Platform
-----------    ------  -------------------------  ----------------
----------------
node1          /cdp
               e0a     cs1                                0/1          N9K-
C9336C
               e0b     cs2                                0/1          N9K-
C9336C

node2          /cdp
               e0a     cs1                                0/2          N9K-
C9336C
               e0b     cs2                                0/2          N9K-
C9336C
```

17. Delete the temporary ISL between cs1 and c1.

```
cs1(config)# no interface port-channel 101
cs1(config)# interface e1/33-34
cs1(config-if-range)# lldp transmit
cs1(config-if-range)# lldp receive
cs1(config-if-range)# no switchport mode trunk
cs1(config-if-range)# no channel-group
cs1(config-if-range)# description 10GbE Node Port
cs1(config-if-range)# spanning-tree bpduguard enable
cs1(config-if-range)# exit
cs1(config)# exit
```

18. Verify the final configuration of the cluster: `network port show -ipspace Cluster`

    Each port should display up for `Link` and healthy for `Health Status`.

```
cluster1::*> network port show -ipspace Cluster

Node: node1

Ignore
                                                 Speed(Mbps) Health
Health
Port        IPspace        Broadcast Domain Link MTU  Admin/Oper  Status
Status
---------   ------------   ---------------- ---- ---- ----------- --------
```

```
------
e0a        Cluster       Cluster           up    9000  auto/10000 healthy
false
e0b        Cluster       Cluster           up    9000  auto/10000 healthy
false


Node: node2

Ignore
                                                      Speed(Mbps) Health
Health
Port       IPspace       Broadcast Domain Link MTU  Admin/Oper  Status
Status
---------  ------------  ---------------- ---- ---- ----------- --------
------
e0a        Cluster       Cluster           up    9000  auto/10000 healthy
false
e0b        Cluster       Cluster           up    9000  auto/10000 healthy
false


4 entries were displayed.


cluster1::*> network interface show -vserver Cluster

           Logical    Status     Network             Current
Current Is
Vserver    Interface  Admin/Oper Address/Mask        Node           Port
Home
---------- ---------- ---------- ----------------- -------------
------- ----
Cluster
           node1_clus1 up/up     169.254.209.69/16  node1          e0a
true
           node1_clus2 up/up     169.254.49.125/16  node1          e0b
true
           node2_clus1 up/up     169.254.47.194/16  node2          e0a
true
           node2_clus2 up/up     169.254.19.183/16  node2          e0b
true


4 entries were displayed.


cluster1::*> network device-discovery show -protocol cdp
```

```
Node/        Local  Discovered
Protocol     Port   Device (LLDP: ChassisID)  Interface       Platform
-----------  ------ ------------------------- ----------------
----------------
node2        /cdp
             e0a    cs1                        0/2             N9K-
C9336C
             e0b    cs2                        0/2             N9K-
C9336C
node1        /cdp
             e0a    cs1                        0/1             N9K-
C9336C
             e0b    cs2                        0/1             N9K-
C9336C

4 entries were displayed.
```

19. Verify that both nodes each have one connection to each switch: `show cdp neighbors`

The following example shows the appropriate results for both switches:

```
cs1# show cdp neighbors

Capability Codes: R - Router, T - Trans-Bridge, B - Source-Route-Bridge
                  S - Switch, H - Host, I - IGMP, r - Repeater,
                  V - VoIP-Phone, D - Remotely-Managed-Device,
                  s - Supports-STP-Dispute

Device-ID          Local Intrfce  Hldtme Capability  Platform      Port
ID
node1              Eth1/1         124    H           FAS2750       e0a
node2              Eth1/2         124    H           FAS2750       e0a
cs2                Eth1/35        179    R S I s     N9K-C9336C
Eth1/35
cs2                Eth1/36        179    R S I s     N9K-C9336C
Eth1/36


cs2# show cdp neighbors

Capability Codes: R - Router, T - Trans-Bridge, B - Source-Route-Bridge
                  S - Switch, H - Host, I - IGMP, r - Repeater,
                  V - VoIP-Phone, D - Remotely-Managed-Device,
                  s - Supports-STP-Dispute

Device-ID          Local Intrfce  Hldtme Capability  Platform      Port
ID
node1              Eth1/1         124    H           FAS2750       e0b
node2              Eth1/2         124    H           FAS2750       e0b
cs1                Eth1/35        179    R S I s     N9K-C9336C
Eth1/35
cs1                Eth1/36        179    R S I s     N9K-C9336C
Eth1/36

Total entries displayed: 4
```

20. Ensure that the cluster network has full connectivity: `cluster ping-cluster -node node-name`

```
cluster1::*> set -priv advanced

Warning: These advanced commands are potentially dangerous; use them
only when
          directed to do so by NetApp personnel.
Do you want to continue? {y|n}: y

cluster1::*> cluster ping-cluster -node node2
Host is node2
Getting addresses from network interface table...
Cluster node1_clus1 169.254.209.69 node1     e0a
Cluster node1_clus2 169.254.49.125 node1     e0b
Cluster node2_clus1 169.254.47.194 node2     e0a
Cluster node2_clus2 169.254.19.183 node2     e0b
Local = 169.254.47.194 169.254.19.183
Remote = 169.254.209.69 169.254.49.125
Cluster Vserver Id = 4294967293
Ping status:
....
Basic connectivity succeeds on 4 path(s)
Basic connectivity fails on 0 path(s)
................
Detected 9000 byte MTU on 4 path(s):
    Local 169.254.19.183 to Remote 169.254.209.69
    Local 169.254.19.183 to Remote 169.254.49.125
    Local 169.254.47.194 to Remote 169.254.209.69
    Local 169.254.47.194 to Remote 169.254.49.125
Larger than PMTU communication succeeds on 4 path(s)
RPC status:
2 paths up, 0 paths down (tcp check)
2 paths up, 0 paths down (udp check)


cluster1::*> set -privilege admin
cluster1::*>
```

21. For ONTAP 9.8 and later, enable the Ethernet switch health monitor log collection feature for collecting switch-related log files, using the following two commands: `system switch ethernet log setup-password` and `system switch ethernet log enable-collection`

    Enter: `system switch ethernet log setup-password`

```
cluster1::*> system switch ethernet log setup-password
Enter the switch name: <return>
The switch name entered is not recognized.
Choose from the following list:
cs1
cs2

cluster1::*> system switch ethernet log setup-password

Enter the switch name: cs1
RSA key fingerprint is e5:8b:c6:dc:e2:18:18:09:36:63:d9:63:dd:03:d9:cc
Do you want to continue? {y|n}::[n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster1::*> system switch ethernet log setup-password

Enter the switch name: cs2
RSA key fingerprint is 57:49:86:a1:b9:80:6a:61:9a:86:8e:3c:e3:b7:1f:b1
Do you want to continue? {y|n}:: [n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>
```

Followed by: `system switch ethernet log enable-collection`

```
cluster1::*> system  switch ethernet log enable-collection

Do you want to enable cluster log collection for all nodes in the
cluster?
{y|n}: [n] y

Enabling cluster switch log collection.

cluster1::*>
```

(i) | If any of these commands return an error, contact NetApp support.

22. For ONTAP releases 9.5P16, 9.6P12, and 9.7P10 and later patch releases, enable the Ethernet switch health monitor log collection feature for collecting switch-related log files, using the commands: `system cluster-switch log setup-password` and `system cluster-switch log enable-collection`

Enter: `system cluster-switch log setup-password`

```
cluster1::*> system cluster-switch log setup-password
Enter the switch name: <return>
The switch name entered is not recognized.
Choose from the following list:
cs1
cs2

cluster1::*> system cluster-switch log setup-password

Enter the switch name: cs1
RSA key fingerprint is e5:8b:c6:dc:e2:18:18:09:36:63:d9:63:dd:03:d9:cc
Do you want to continue? {y|n}::[n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster1::*> system cluster-switch log setup-password

Enter the switch name: cs2
RSA key fingerprint is 57:49:86:a1:b9:80:6a:61:9a:86:8e:3c:e3:b7:1f:b1
Do you want to continue? {y|n}:: [n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>
```

Followed by: `system cluster-switch log enable-collection`

```
cluster1::*> system cluster-switch log enable-collection

Do you want to enable cluster log collection for all nodes in the
cluster?
{y|n}: [n] y

Enabling cluster switch log collection.

cluster1::*>
```

(i) | If any of these commands return an error, contact NetApp support.

23. If you suppressed automatic case creation, reenable it by invoking an AutoSupport message: `system node autosupport invoke -node * -type all -message MAINT=END`