



Migrate from CN1610 switches to BES-53248 switches

Cluster and storage switches

NetApp
September 12, 2022

This PDF was generated from <https://docs.netapp.com/us-en/ontap-systems-switches/switch-bes-53248/migrate-cn1610-overview.html> on September 12, 2022. Always check docs.netapp.com for the latest.

Table of Contents

- Migrate from CN1610 switches to BES-53248 switches 1
 - Migrate CN1610 cluster switches to Broadcom-supported BES-53248 cluster switches 1
 - How to migrate CN1610 cluster switches to BES-53248 cluster switches 1

Migrate from CN1610 switches to BES-53248 switches

Migrate CN1610 cluster switches to Broadcom-supported BES-53248 cluster switches

You must be aware of certain configuration information, port connections, and cabling requirements when you migrate CN1610 cluster switches to Broadcom-supported BES-53248 cluster switches.

- The following cluster switches are supported:
 - CN1610
 - BES-53248
- The cluster switches support the following node connections:
 - NetApp CN1610: ports 0/1 through 0/12 (10 GbE)
 - BES-53248: ports 0/1-0/16 (10/25 GbE)



Additional ports can be activated by purchasing port licenses.

- The cluster switches use the following inter-switch link (ISL) ports:
 - NetApp CN1610: ports 0/13 through 0/16 (10 GbE)
 - BES-53248: ports 0/55-0/56 (100 GbE)
- The [NetApp Hardware Universe](#) contains information about ONTAP compatibility, supported EFOS firmware, and cabling to BES-53248 cluster switches.
- The appropriate ISL cabling is as follows:
 - **Beginning:** For CN1610 to CN1610 (SFP+ to SFP+), four SFP+ optical fiber or copper direct-attach cables.
 - **Interim:** For CN1610 to BES-53248 (SFP+ to SFP28), four 10G SFP+ optical transceiver/fiber or copper direct-attach cables.
 - **Final:** For BES-53248 to BES-53248 (QSFP28 to QSFP28), two QSFP28 optical transceivers/fiber or copper direct-attach cables.



After your migration completes, you might need to install the required configuration file to support the Cluster Switch Health Monitor (CSHM) for BES-53248 cluster switches.

See [Install the Cluster Switch Health Monitor \(CSHM\) configuration file](#) and [Configure the cluster switch log collection feature](#) for the steps required to enable cluster health switch log collection used for collecting switch-related log files.

How to migrate CN1610 cluster switches to BES-53248 cluster switches

To replace the existing CN1610 cluster switches in a cluster with Broadcom-supported

BES-53248 cluster switches, you must perform a specific sequence of tasks.

What you'll need

The examples in this procedure use two nodes, each deploying two 10 GbE cluster interconnect ports: e0a and e0b.

The examples in this procedure use the following switch and node nomenclature:

- The command outputs might vary depending on different releases of ONTAP software.
- The CN1610 switches to be replaced are CL1 and CL2.
- The BES-53248 switches to replace the CN1610 switches are cs1 and cs2.
- The nodes are node1 and node2.
- The switch CL2 is replaced by cs2 first, followed with CL1 by cs1.
- The BES-53248 switches are pre-loaded with the supported versions of Reference Configuration File (RCF) and Ethernet Fabric OS (EFOS) with ISL cables connected on ports 55 and 56.
- The cluster LIF names are node1_clus1 and node1_clus2 for node1, and node2_clus1 and node2_clus2 for node2.

About this task

This procedure covers the following scenario:

- The cluster starts with two nodes connected to two CN1610 cluster switches.
- CN1610 switch CL2 is replaced by BES-53248 switch cs2:
 - Disconnect the cables from all cluster ports on all nodes connected to CL2, and then use supported cables to reconnect the ports to the new cluster switch cs2.
 - Disconnect the cables between ISL ports CL1 and CL2, and then use supported cables to reconnect the ports from CL1 to cs2.
- CN1610 switch CL1 is replaced by BES-53248 switch cs1:
 - Disconnect the cables from all cluster ports on all nodes connected to CL1, and then use supported cables to reconnect the ports to the new cluster switch cs1.
 - Disconnect the cables between ISL ports CL1 and cs2, and then use supported cables to reconnect the ports from cs1 to cs2.

Steps

1. If AutoSupport is enabled on this cluster, suppress automatic case creation by invoking an AutoSupport message:

```
system node autosupport invoke -node * -type all -message MAINT=xh
```

where x is the duration of the maintenance window in hours.



The AutoSupport message notifies technical support of this maintenance task so that automatic case creation is suppressed during the maintenance window.

The following command suppresses automatic case creation for two hours:

```
cluster1::*> system node autosupport invoke -node \* -type all -message  
MAINT=2h
```

2. Change the privilege level to advanced, entering **y** when prompted to continue:

```
set -privilege advanced
```

The advanced prompt (*>) appears.

3. Verify that auto-revert is enabled on all cluster LIFs:

```
network interface show -vserver Cluster -fields auto-revert
```

```
cluster1::*> network interface show -vserver Cluster -fields auto-revert
```

| Vserver | Logical Interface | Auto-revert |
|---------|----------------------|-------------|
| Cluster | node1_clus1 | true |
| | node1_clus2 | true |
| | node2_clus1 | true |
| | node2_clus2 | true |

4. Display information about the devices in your configuration:

```
network device-discovery show -protocol cdp
```

The following example displays how many cluster interconnect interfaces have been configured in each node for each cluster interconnect switch:

```
cluster1::*> network device-discovery show -protocol cdp
```

| Node/ Protocol | Local Port | Discovered Device (LLDP: ChassisID) | Interface | Platform |
|-------------------|---------------|--|-----------|----------|
| node2 | /cdp | | | |
| | e0a | CL2 | 0/2 | CN1610 |
| | e0b | CL1 | 0/2 | CN1610 |
| node1 | /cdp | | | |
| | e0a | CL2 | 0/1 | CN1610 |
| | e0b | CL1 | 0/1 | CN1610 |

5. Determine the administrative or operational status for each cluster interface.

- a. Display the cluster network port attributes:

```
network port show -ipspace Cluster
```

```
cluster1::*> network port show -ipspace Cluster
```

```
Node: node1
```

```
Ignore
```

| | | | | | | Speed(Mbps) | Health |
|--------|---------|-----------|--------|------|------|-------------|---------|
| Health | | | | | | | |
| Port | IPspace | Broadcast | Domain | Link | MTU | Admin/Oper | Status |
| Status | | | | | | | |
| ----- | ----- | ----- | ----- | ---- | ---- | ----- | |
| ----- | ----- | | | | | | |
| e0a | Cluster | Cluster | | up | 9000 | auto/10000 | healthy |
| false | | | | | | | |
| e0b | Cluster | Cluster | | up | 9000 | auto/10000 | healthy |
| false | | | | | | | |

```
Node: node2
```

```
Ignore
```

| | | | | | | Speed(Mbps) | Health |
|--------|---------|-----------|--------|------|------|-------------|---------|
| Health | | | | | | | |
| Port | IPspace | Broadcast | Domain | Link | MTU | Admin/Oper | Status |
| Status | | | | | | | |
| ----- | ----- | ----- | ----- | ---- | ---- | ----- | |
| ----- | ----- | | | | | | |
| e0a | Cluster | Cluster | | up | 9000 | auto/10000 | healthy |
| false | | | | | | | |
| e0b | Cluster | Cluster | | up | 9000 | auto/10000 | healthy |
| false | | | | | | | |

b. Display information about the logical interfaces:

```
network interface show -vserver Cluster
```

```
cluster1::*> network interface show -vserver Cluster
```

| | Logical | Status | Network | Current |
|------------|-------------|------------|-------------------|---------|
| Current Is | | | | |
| Vserver | Interface | Admin/Oper | Address/Mask | Node |
| Port | Home | | | |
| ----- | | | | |
| ----- | | | | |
| Cluster | | | | |
| e0a | node1_clus1 | up/up | 169.254.209.69/16 | node1 |
| e0b | node1_clus2 | up/up | 169.254.49.125/16 | node1 |
| e0a | node2_clus1 | up/up | 169.254.47.194/16 | node2 |
| e0b | node2_clus2 | up/up | 169.254.19.183/16 | node2 |
| | | | | |

- Verify that the appropriate port licenses, RCF, and EFOS image are installed on the new BES-53248 switches as necessary for your requirements, and make any essential site customizations, such as users and passwords, network addresses, and so on.
- Ping the remote cluster interfaces:

```
cluster ping-cluster -node node-name
```

The following example shows how to ping the remote cluster interfaces:

```

cluster1::*> cluster ping-cluster -node node2

Host is node2
Getting addresses from network interface table...
Cluster node1_clus1 169.254.209.69  node1      e0a
Cluster node1_clus2 169.254.49.125  node1      e0b
Cluster node2_clus1 169.254.47.194  node2      e0a
Cluster node2_clus2 169.254.19.183  node2      e0b

Local = 169.254.47.194 169.254.19.183
Remote = 169.254.209.69 169.254.49.125
Cluster Vserver Id = 4294967293
Ping status:

Basic connectivity succeeds on 4 path(s)
Basic connectivity fails on 0 path(s)

Detected 9000 byte MTU on 4 path(s):
    Local 169.254.47.194 to Remote 169.254.209.69
    Local 169.254.47.194 to Remote 169.254.49.125
    Local 169.254.19.183 to Remote 169.254.209.69
    Local 169.254.19.183 to Remote 169.254.49.125
Larger than PMTU communication succeeds on 4 path(s)

RPC status:
2 paths up, 0 paths down (tcp check)
2 paths up, 0 paths down (udp check)

```

8. Shut down the ISL ports 13 through 16 on the active CN1610 switch CL1:

shutdown

The following example shows how to shut down ISL ports 13 through 16 on the CN1610 switch CL1:

```

(CL1)# configure
(CL1)(Config)# interface 0/13-0/16
(CL1)(Interface 0/13-0/16)# shutdown
(CL1)(Interface 0/13-0/16)# exit
(CL1)(Config)# exit
(CL1)#

```

9. Build a temporary ISL between CN1610 CL1 and new BES-53248 cs2. The ISL will only be defined on cs2 as the existing ISL on CL1 can be reused.

The following example builds a temporary ISL on cs2 (ports 13-16) to be connected to the existing ISL on CL1 (ports 13-16):


```

(cs2)# configure
(cs2) (Config)# port-channel name 1/2 temp-isl-cn1610
(cs2) (Config)# interface 0/13-0/16
(cs2) (Interface 0/13-0/16)# no spanning-tree edgeport
(cs2) (Interface 0/13-0/16)# addport 1/2
(cs2) (Interface 0/13-0/16)# exit
(cs2) (Config)# interface lag 2
(cs2) (Interface lag 2)# mtu 9216
(cs2) (Interface lag 2)# port-channel load-balance 7
(cs2) (Config)# exit

(cs2)# show port-channel 1/2
Local Interface..... 1/2
Channel Name..... temp-isl-cn1610
Link State..... Down
Admin Mode..... Enabled
Type..... Static
Port-channel Min-links..... 1
Load Balance Option..... 7
(Enhanced hashing mode)

Mbr      Device/      Port      Port
Ports    Timeout      Speed     Active
-----
0/13      actor/long    10G Full  False
          partner/long
0/14      actor/long    10G Full  False
          partner/long
0/15      actor/long    10G Full  False
          partner/long
0/16      actor/long    10G Full  False
          partner/long

```

10. On all nodes, remove the cables that are attached to the CN1610 switch CL2.

You must then reconnect the disconnected ports on all nodes to the new BES-53248 switch cs2. Refer to the [NetApp Hardware Universe](#) for approved cabling options.

11. Remove four ISL cables from ports 13 to 16 on the CN1610 switch CL2.

You must attach appropriate approved cabling connecting port 0/13 to 0/16 on the new BES-53248 switch cs2, to ports 13 to 16 on the existing CN1610 switch CL1.

12. Bring up ISLs 13 through 16 on the active CN1610 switch CL1.

The following example illustrates the process of bringing up ISL ports 13 through 16 on CL1:

```
(CL1)# configure
(CL1) (Config)# interface 0/13-0/16
(CL1) (Interface 0/13-0/16,3/1)# no shutdown
(CL1) (Interface 0/13-0/16,3/1)# exit
(CL1) (Config)# exit
(CL1)#
```

13. Verify that the ISLs are **up** on the CN1610 switch CL1:

```
show port-channel
```

The Link State should be Up, Type should be Static, and Port Active should be True for ports 0/13 to 0/16:

```
(CL2)# show port-channel 3/1
Local Interface..... 3/1
Channel Name..... ISL-LAG
Link State..... Up
Admin Mode..... Enabled
Type..... Static
Load Balance Option..... 7
```

(Enhanced hashing mode)

| Mbr Ports | Device/ Timeout | Port Speed | Port Active |
|--------------|----------------------------|---------------|----------------|
| ----- | ----- | ----- | ----- |
| 0/13 | actor/long partner/long | 10 Gb Full | True |
| 0/14 | actor/long partner/long | 10 Gb Full | True |
| 0/15 | actor/long partner/long | 10 Gb Full | True |
| 0/16 | actor/long partner/long | 10 Gb Full | True |

14. Verify that the ISL ports are up on the BES-53248 switch:

```
show port-channel
```

```
(cs2)# show port-channel 1/2
```

```
Local Interface..... 1/2
Channel Name..... temp-isl-cn1610
Link State..... Up
Admin Mode..... Enabled
Type..... Static
Port-channel Min-links..... 1
Load Balance Option..... 7
```

```
(Src/Dest MAC, VLAN, EType, incoming port)
```

| Mbr Ports | Device/ Timeout | Port Speed | Port Active |
|--------------|----------------------------|---------------|----------------|
| 0/13 | actor/long partner/long | 10G Full | True |
| 0/14 | actor/long partner/long | 10G Full | True |
| 0/15 | actor/long partner/long | 10G Full | True |
| 0/16 | actor/long partner/long | 10G Full | True |

15. Verify that all of the cluster interconnect ports are reverted to their home ports:

```
network interface show -vserver Cluster
```

```
cluster1::*> network interface show -vserver Cluster
```

| Logical | Status | Network | Current | |
|-------------|-----------|-------------------|--------------|------|
| Current Is | | | | |
| Vserver | Interface | Admin/Oper | Address/Mask | Node |
| Port | Home | | | |
| Cluster | | | | |
| node1_clus1 | up/up | 169.254.209.69/16 | node1 | e0a |
| node1_clus2 | up/up | 169.254.49.125/16 | node1 | e0b |
| node2_clus1 | up/up | 169.254.47.194/16 | node2 | e0a |
| node2_clus2 | up/up | 169.254.19.183/16 | node2 | e0b |

16. Verify that all of the cluster ports are connected:

```
network port show -ipspace Cluster
```

The following example shows the result of the previous command, verifying that all of the cluster interconnects are up:

```
cluster1::*> network port show -ipspace Cluster
```

Node: node1

Ignore

| | | | | | | Speed(Mbps) | Health |
|--------|---------|-----------|--------|------|------|-------------|---------|
| Health | | | | | | | |
| Port | IPspace | Broadcast | Domain | Link | MTU | Admin/Oper | Status |
| Status | | | | | | | |
| ----- | ----- | ----- | ----- | ---- | ---- | ----- | ----- |
| ----- | | | | | | | |
| e0a | Cluster | Cluster | | up | 9000 | auto/10000 | healthy |
| false | | | | | | | |
| e0b | Cluster | Cluster | | up | 9000 | auto/10000 | healthy |
| false | | | | | | | |

Node: node2

Ignore

| | | | | | | Speed(Mbps) | Health |
|--------|---------|-----------|--------|------|------|-------------|---------|
| Health | | | | | | | |
| Port | IPspace | Broadcast | Domain | Link | MTU | Admin/Oper | Status |
| Status | | | | | | | |
| ----- | ----- | ----- | ----- | ---- | ---- | ----- | ----- |
| ----- | | | | | | | |
| e0a | Cluster | Cluster | | up | 9000 | auto/10000 | healthy |
| false | | | | | | | |
| e0b | Cluster | Cluster | | up | 9000 | auto/10000 | healthy |
| false | | | | | | | |

17. Ping the remote cluster interfaces:

```
cluster ping-cluster -node node-name
```

The following example shows how to ping the remote cluster interfaces:

```

cluster1::*> cluster ping-cluster -node node2
Host is node2
Getting addresses from network interface table...
Cluster node1_clus1 169.254.209.69  node1      e0a
Cluster node1_clus2 169.254.49.125  node1      e0b
Cluster node2_clus1 169.254.47.194  node2      e0a
Cluster node2_clus2 169.254.19.183  node2      e0b
Local = 169.254.47.194 169.254.19.183
Remote = 169.254.209.69 169.254.49.125
Cluster Vserver Id = 4294967293
Ping status:
....
Basic connectivity succeeds on 4 path(s)
Basic connectivity fails on 0 path(s)
.....
Detected 9000 byte MTU on 4 path(s):
    Local 169.254.47.194 to Remote 169.254.209.69
    Local 169.254.47.194 to Remote 169.254.49.125
    Local 169.254.19.183 to Remote 169.254.209.69
    Local 169.254.19.183 to Remote 169.254.49.125
Larger than PMTU communication succeeds on 4 path(s)
RPC status:
2 paths up, 0 paths down (tcp check)
2 paths up, 0 paths down (udp check)

```

18. On all nodes, remove the cables that are attached to the CN1610 switch CL1.

You must then reconnect the disconnected ports on all nodes to the new BES-53248 switch cs1. Refer to the [NetApp Hardware Universe](#) for approved cabling options.

19. Remove four ISL cables from ports 13 to 16 on BES-53248 switch cs2.
20. Remove the temporary port-channel 2 on cs2.

The following example removes port-channel 2 and copies the running-configuration file to the startup-configuration file:

```
(cs2)# configure
(cs2) (Config)# deleteport 1/2 all
(cs2) (Config)# interface 0/13-0/16
(cs2) (Interface 0/13-0/16)# spanning-tree edgeport
(cs2) (Interface 0/13-0/16)# exit
(cs2) (Config)# exit
(cs2)# write memory
```

This operation may take a few minutes.

Management interfaces will not be available during this time.

Are you sure you want to save? (y/n) **y**

Config file 'startup-config' created successfully .

21. Verify the status of the cluster node port:

```
network port show -ipspace Cluster
```

The following example verifies that all of the cluster interconnect ports on node1 and node2 are up:

```
cluster1::*> network port show -ipspace Cluster
```

Node: node1

Ignore

| | | | | | | Speed(Mbps) | Health |
|--------|---------|-----------|--------|------|------|-------------|---------|
| Health | | | | | | | |
| Port | IPspace | Broadcast | Domain | Link | MTU | Admin/Oper | Status |
| Status | | | | | | | |
| ----- | ----- | ----- | ----- | ---- | ---- | ----- | ----- |
| e0a | Cluster | Cluster | | up | 9000 | auto/10000 | healthy |
| false | | | | | | | |
| e0b | Cluster | Cluster | | up | 9000 | auto/10000 | healthy |
| false | | | | | | | |

Node: node2

Ignore

| | | | | | | Speed(Mbps) | Health |
|--------|---------|-----------|--------|------|------|-------------|---------|
| Health | | | | | | | |
| Port | IPspace | Broadcast | Domain | Link | MTU | Admin/Oper | Status |
| Status | | | | | | | |
| ----- | ----- | ----- | ----- | ---- | ---- | ----- | ----- |
| e0a | Cluster | Cluster | | up | 9000 | auto/10000 | healthy |
| false | | | | | | | |
| e0b | Cluster | Cluster | | up | 9000 | auto/10000 | healthy |
| false | | | | | | | |

22. Verify that the interface is now home:

```
network interface show -vserver Cluster
```

The following example shows the status of cluster interconnect interfaces are up and Is home for node1 and node2:

```
cluster1::*> network interface show -vserver Cluster
```

| | Logical | Status | Network | Current | Current |
|---------|-------------|------------|-------------------|---------|---------|
| Is | | | | | |
| Vserver | Interface | Admin/Oper | Address/Mask | Node | Port |
| Home | | | | | |
| ----- | ----- | ----- | ----- | ----- | ----- |
| Cluster | | | | | |
| true | node1_clus1 | up/up | 169.254.209.69/16 | node1 | e0a |
| true | node1_clus2 | up/up | 169.254.49.125/16 | node1 | e0b |
| true | node2_clus1 | up/up | 169.254.47.194/16 | node2 | e0a |
| true | node2_clus2 | up/up | 169.254.19.183/16 | node2 | e0b |
| true | | | | | |

23. Ping the remote cluster interfaces and then perform a remote procedure call server check:

```
cluster ping-cluster -node node-name
```

The following example shows how to ping the remote cluster interfaces:


```

cluster1::*> cluster ping-cluster -node node2
Host is node2
Getting addresses from network interface table...
Cluster node1_clus1 169.254.209.69 node1 e0a
Cluster node1_clus2 169.254.49.125 node1 e0b
Cluster node2_clus1 169.254.47.194 node2 e0a
Cluster node2_clus2 169.254.19.183 node2 e0b
Local = 169.254.47.194 169.254.19.183
Remote = 169.254.209.69 169.254.49.125
Cluster Vserver Id = 4294967293
Ping status:

Basic connectivity succeeds on 4 path(s)
Basic connectivity fails on 0 path(s)
.....
Detected 9000 byte MTU on 4 path(s):
    Local 169.254.47.194 to Remote 169.254.209.69
    Local 169.254.47.194 to Remote 169.254.49.125
    Local 169.254.19.183 to Remote 169.254.209.69
    Local 169.254.19.183 to Remote 169.254.49.125
Larger than PMTU communication succeeds on 4 path(s)
RPC status:
2 paths up, 0 paths down (tcp check)
2 paths up, 0 paths down (udp check)

```

24. Display the information about the devices in your configuration:

```
network device-discovery show -protocol cdp
```

The following examples show node1 and node2 have been migrated from CN1610 CL2 and CL1 to BES-53248 cs2 and cs1:

```

cluster1::*> network device-discovery show -protocol cdp
Node/      Local  Discovered
Protocol   Port   Device (LLDP: ChassisID)  Interface      Platform
-----
node1      /cdp
           e0a    cs2                      0/1            BES-53248
           e0b    cs1                      0/1            BES-53248
node2      /cdp
           e0a    cs2                      0/2            BES-53248
           e0b    cs1                      0/2            BES-53248

```

25. Remove the replaced CN1610 switches if they are not automatically removed:

```
system cluster-switch delete -device device-name
```

The following example shows how to remove the CN1610 switches:

```
cluster::*> system cluster-switch delete -device CL2  
cluster::*> system cluster-switch delete -device CL1
```

26. If you suppressed automatic case creation, re-enable it by invoking an AutoSupport message:

```
system node autosupport invoke -node * -type all -message MAINT=END
```

```
cluster::*> system node autosupport invoke -node \ -type all -message  
MAINT=END*
```

After you finish

See [Install the Cluster Switch Health Monitor \(CSHM\) configuration file](#) and [Configure the cluster switch log collection feature](#) for the steps required to enable cluster health switch log collection used for collecting switch-related log files.

Related information

[Hardware Universe](#)

[Broadcom-supported BES-53248 switches setup and configuration](#)

Copyright Information

Copyright © 2022 NetApp, Inc. All rights reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means-graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system- without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP "AS IS" AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark Information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.